

Population Based Training of Neural Networks

Max Jaderberg, Valentin Dalibard, Simon Osindero

DeepMind

arXiv:1711.09846

Information & Intelligence System Lab.

Sang Heon Lee

lawlee1@naver.com


2019/01/11

Introduction

- Hyperparameter of Model
 - Hyperparameters must be properly tuned to fully unlock network performance
 - Tuning process is computationally expensive → Hyperparameter search methods
- Hyperparameter search methods
 - Parallel search
 - Training models with diff. hyperparameters in parallel, select the best one
 - Require a lot of computational resources
 - Sequential optimization
 - Using information obtained from earlier training runs
 - Require a lot of time

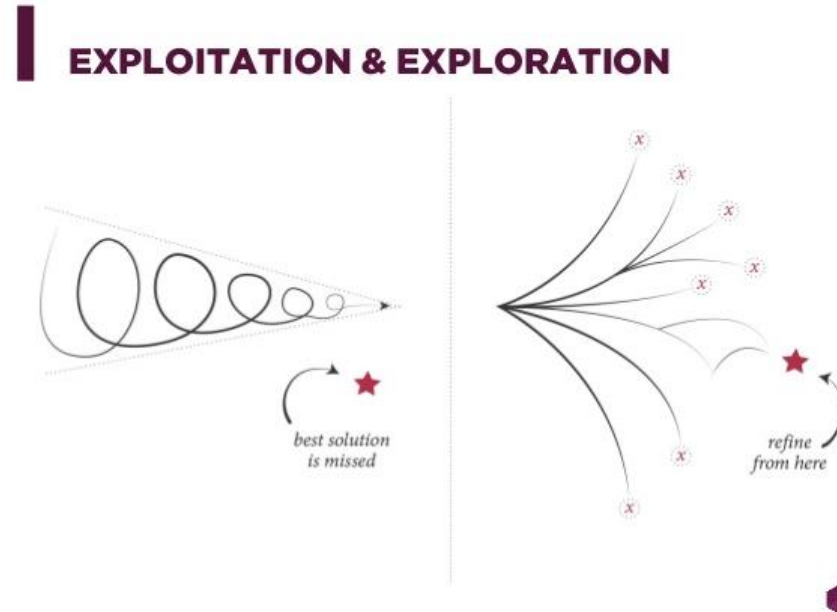
Introduction (Cont'd)

- Present simple hyperparameter searching methods
 - Bridges between parallel search and sequential optimization methods
 - *Population Based Training (PBT)*
 - Based on genetic algorithm
 - Information sharing across a population
 - Transfer **parameters** and **hyperparameters** between members of the population




Related Works

- Optimization : Exploitation vs. Exploration
 - Exploitation : choose the best point until now
 - Exploration : choose a new point to get more information
 - Trade-off relations : the proper control of these two behaviors is core



Related Works (Cont'd)

- Sequential optimization : Bayesian Optimization
 - Utilize the information of previous trial points
 - Goal is finding $\mathbf{x}^* = \arg \max_{\mathbf{x} \in \mathcal{X}} f(\mathbf{x})$
- Approximate expensive function f using a probabilistic surrogate model




Algorithm 1 Bayesian optimization

- 1: **for** $n = 1, 2, \dots$ **do**
 - 2: select new \mathbf{x}_{n+1} by optimizing acquisition function α
$$\mathbf{x}_{n+1} = \arg \max_{\mathbf{x}} \alpha(\mathbf{x}; \mathcal{D}_n)$$
 - 3: query objective function to obtain y_{n+1}
 - 4: augment data $\mathcal{D}_{n+1} = \{\mathcal{D}_n, (\mathbf{x}_{n+1}, y_{n+1})\}$
 - 5: update statistical model
 - 6: **end for**
-

- Require **a lot of time** to find the best solution

Related Works (Cont'd)


- Parallel search : Random search



- Waste computation on bad hyperparameters
- Fails to utilize the **information of history**
- *Hyperband* (Li et al., 2016)
 - Allocate more budget to more promising hyperparameter configurations
 - Has same problem as random search

Proposed Method


- *Population Based Training (PBT)*



- Start with random search
- Allow workers to **share information**
- Workers can **exploit** for model selection, and **explore** new hyperparameters
- Genetic algorithm acting on a timescale which allows gradient based learning

Proposed Method (Cont'd)


- *Population Based Training (PBT)*



- Randomly initialize model weights
- Randomly initialize hyperparameters from a prior distribution

Proposed Method (Cont'd)


- *Population Based Training (PBT)*



- Allow training for enough steps

Proposed Method (Cont'd)

- *Population Based Training (PBT)*




- **Exploit**

- Each workers compares its performance to the population. If bad, abandon it and replace the model and hyperparameter with better worker
- **Binary tournament** – random two, better wins
- **Truncation selection** – if in bottom 20%, replace with top 20%

Proposed Method (Cont'd)

- *Population Based Training (PBT)*




- **Explore**

- Mutate the hyperparameters that were replaced
- **Perturb** – randomly perturbed by a factor e.g. 1.2
- **Resample** – resampled from the initial prior distribution defined

Proposed Method (Cont'd)


- *Population Based Training (PBT)*



- **Step:** perform steps of regular gradient-based training
- **Exploit:** if worker is bad, then replace it with better partial model
- **Explore:** mutate hyperparameters that replaced
- **Repeat**

Proposed Method (Cont'd)

- *Population Based Training (PBT)*



- Combines model optimization and hyperparameter refinement
- Exploit can optimize for non-differentiable & expensive metrics
 - accuracy on test set, BLEU scores, human normalized performance, ...
- All workers benefit from the exploration luck

Experiments

- Apply PBT to 3 diff. learning problems
 - **RL** (Reinforcement Learning), **MT** (Machine Translation)
GAN (Generative Adversarial Networks)
- 1. Deep Reinforcement Learning
 - Training of neural network agent with RL
 - Find a policy π to maximize expected episodic reward $E_{\pi}[R]$
 - 3 Tasks and models
 - DeepMind Lab, *UNREAL* (Jaderberg et al., 2016)
 - Atari games, *Feudal Networks* (Vezhnevets et al., 2017)
 - StarCraft 2, *A3C* baseline agents (Vinyals et al., 2017)

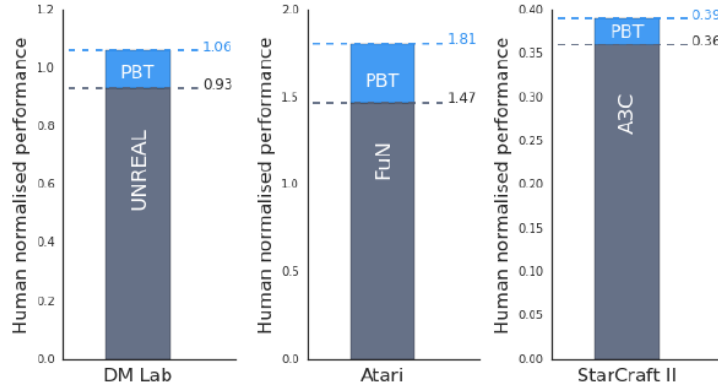
Experiments (Cont'd)

1. Deep Reinforcement Learning

- Hyperparameters
 - learning rate, entropy cost, unroll length, intrinsic reward cost
- Step
 - Step of gradient descent with RMSProp
- Eval
 - Last 10 episodic rewards
- Ready
 - between $10^6 \sim 10^7$ agent steps have elapsed
- Baseline
 - Random search with the same number of workers

Experiments (Cont'd)

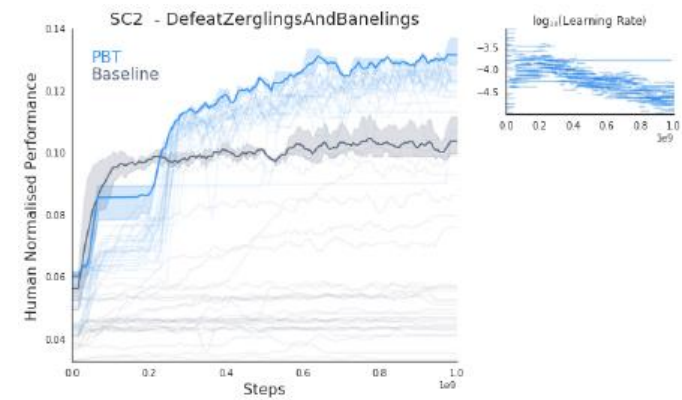
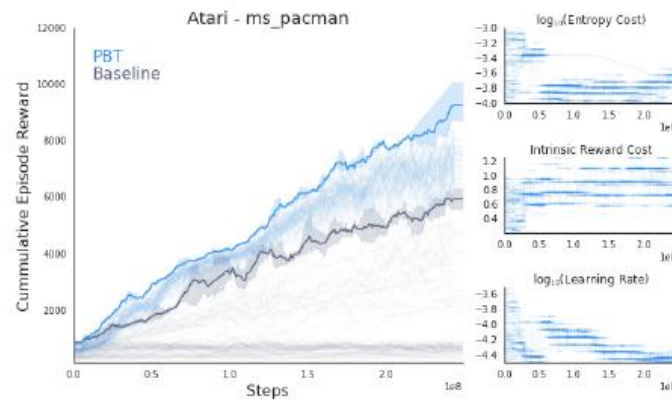
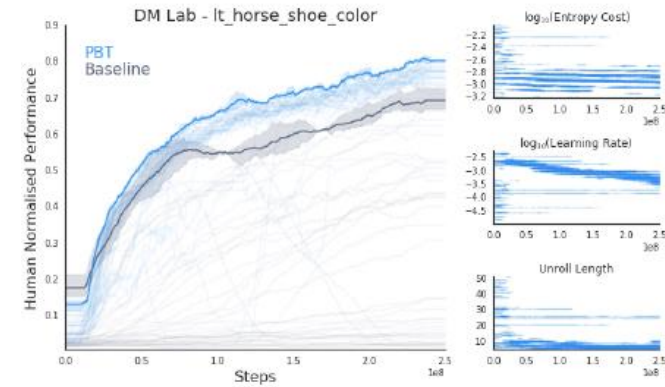
1. Deep Reinforcement Learning



workers : 40

80

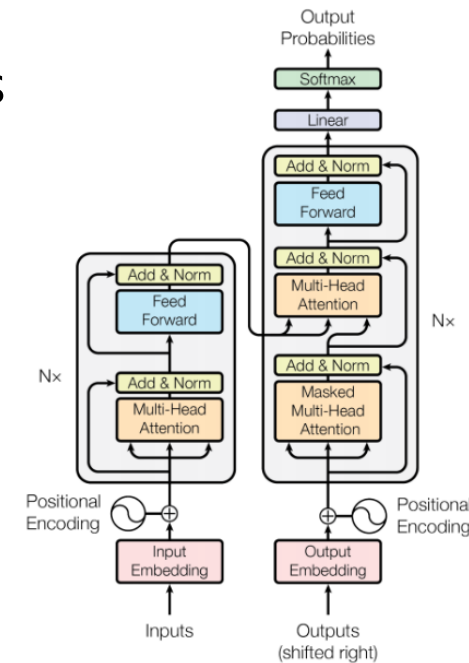
30



Experiments (Cont'd)

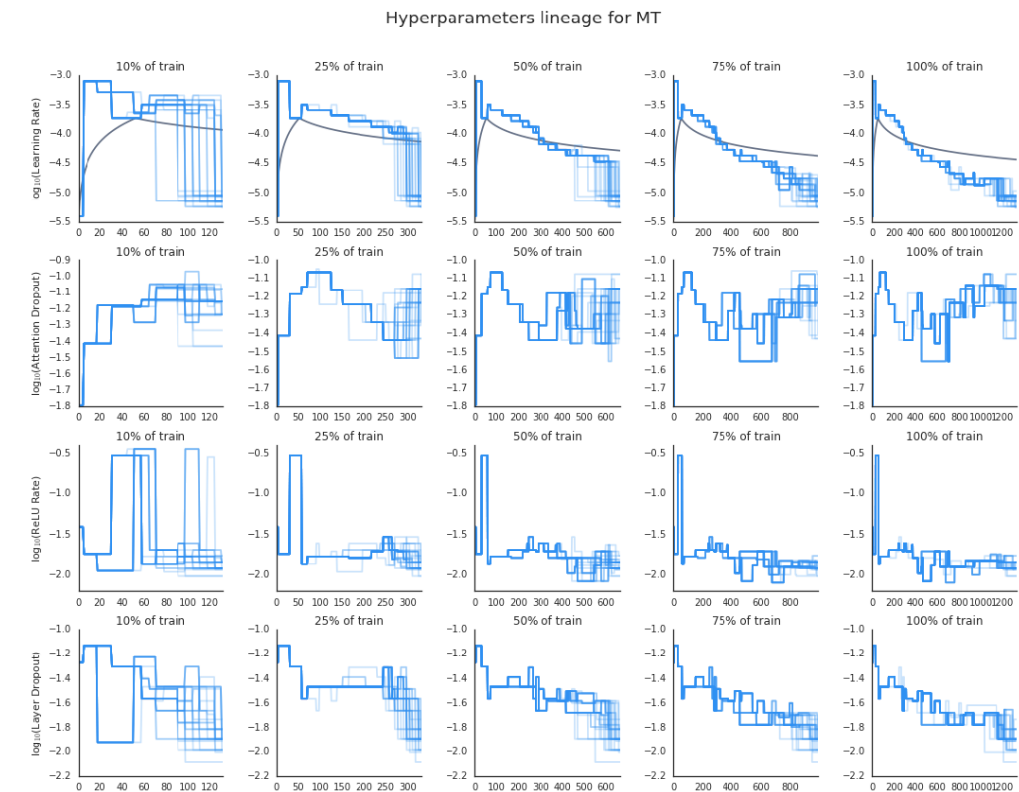


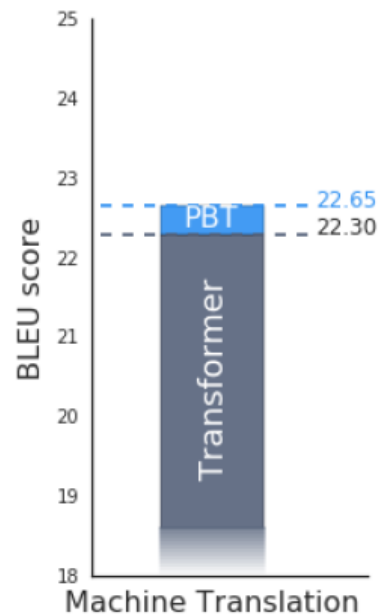
2. Machine Translation

- *Transformer* networks (Vaswani et al., 2017), English to German on WMT 2014
- 32 workers with 400×10^3 steps
- Hyperparameters
 - learning rate, attention dropout, layer dropout, ReLU dropout rates
- Step : step of Adam
- Eval : **BLEU score** on WMT *newstest2012* dataset
- Exploit : **binary tournament**
- Explore : **perturb** (1.2 or 0.8)
- Baseline model : highly optimized hyperparameter values
 - Result of hand tuning and Bayesian Optimization



Experiments (Cont'd)

2. Machine Translation



small Transformer network with small batch size

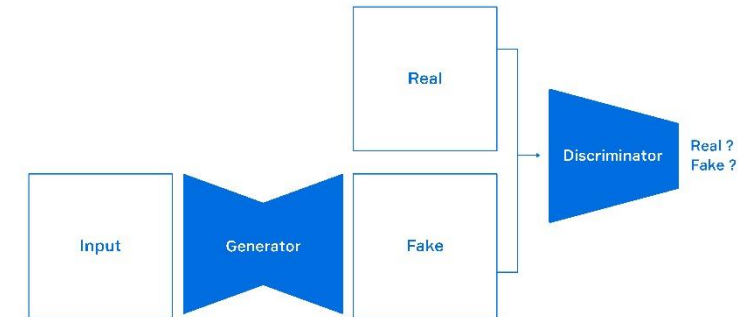
Experiments (Cont'd)

3. Generative Adversarial Networks

- Model : *DCGAN* (Radford et al., 2016) CIFAR trained
- Hyperparameters
 - G's learning rate, D's learning rate

- Eval

- Inception score : $IS(G) = \exp (\mathbb{E}_{\mathbf{x} \sim p_g} D_{KL}(p(y|\mathbf{x}) \parallel p(y)))$
 - Outputs of pretrained CIFAR classifier
- Exploit : **both** (binary tournament, truncation selection)
- Explore : **perturb** (2.0 or 0.5)
- 45 workers
- Baseline model : Best among hand-design annealing strategies




Good G
→ low of $p(y|x)$, high of $p(y)$
→ high of $IS(G)$

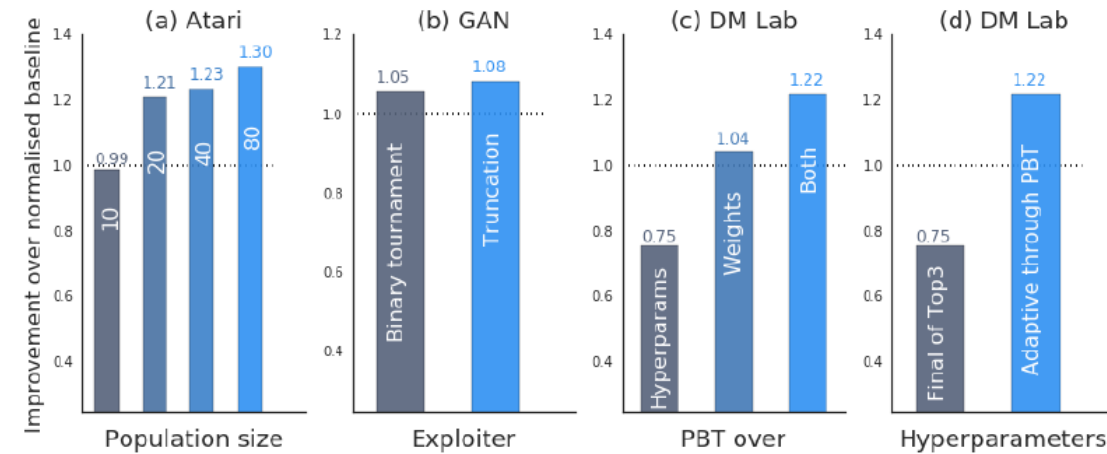
Experiments (Cont'd)

3. Generative Adversarial Networks

GAN results



Ablation study



Conclusion

- Proposed Population Based Training (PBT)
 - Based on Genetic algorithm
 - Bridges between **parallel search** and **sequential optimization** methods
 - Optimize over weights and hyperparameters jointly
 - Improvements in accuracy, stability across a wide range of domains
 - Discovers an **adaptive schedule** rather than fixed set of hyperparameters